

Introduction to Recurrent Neural Network

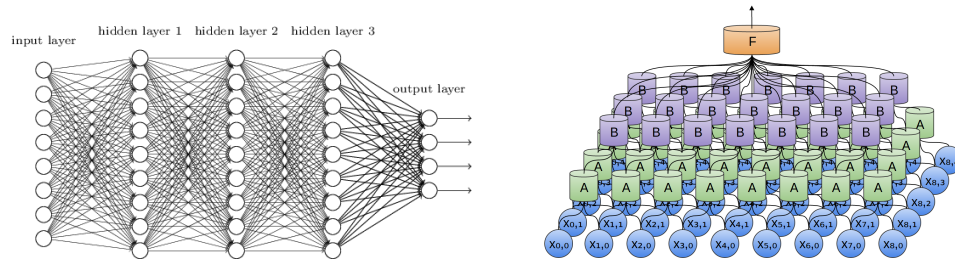
Bokan Bao
Professor Dasgupta
CSE 254

Outlines


- RNN introduction
 - Limitation of Vanilla Neural Network
 - Structure of the RNN
 - PyTorch implementation
- Character-Level prediction
 - Unsmoothed Maximum Likelihood Character Level Language Model
 - Three-layer RNN

Vanilla Neural Network

They accept a fixed-sized vector as input (e.g. an image) and produce a fixed-sized vector as output (e.g. probabilities of different classes).




Recurrent Neural Network





Recurrent Neural Network

- Flexible, a sequence of the data
- Music, text, motion capture
- The predictive distribution depends on the previous inputs.



A sequence of many things



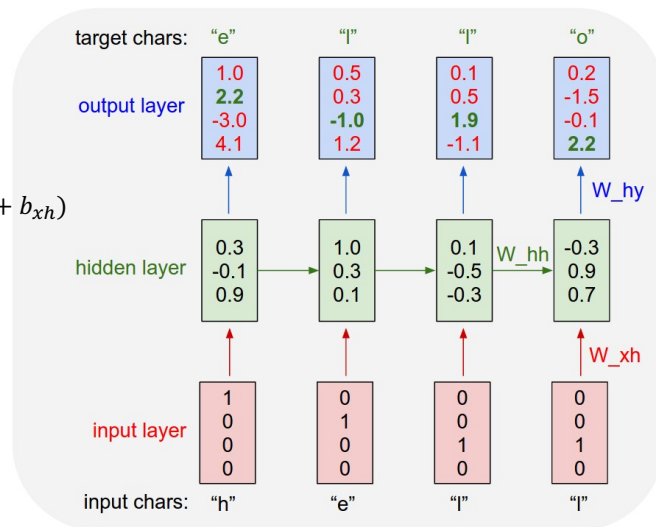

Simple RNN model

$$x_t \rightarrow W_{xh}x_t + b_{xh}$$

$$h_{t-1} \rightarrow h_t$$

$$h_t = \tanh(W_{hh}h_{t-1} + b_{hh} + W_{xh}x_t + b_{xh})$$

$$y_t = W_{hy}h_t + b_{hy}$$

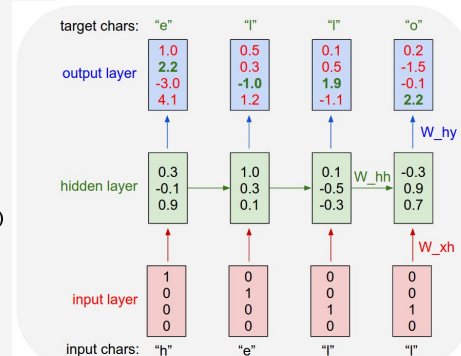


Pytorch implementation

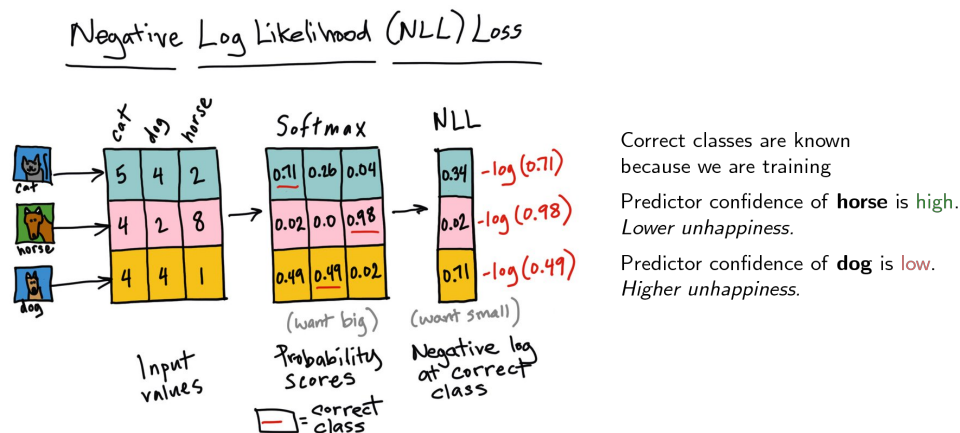
```
class RNN(nn.Module):
    def __init__(self, input_size, hidden_size, output_size):
        super().__init__()
        self.hidden_size = hidden_size
        self.W_xh = nn.Linear(input_size, hidden_size)
        self.W_hh = nn.Linear(hidden_size, hidden_size)
        self.W_hy = nn.Linear(hidden_size, output_size)
        self.act = nn.Tanh()
        self.dropout = nn.Dropout(0.1)
        self.softmax = nn.LogSoftmax(dim=1)

    def forward(self, _input, hidden):
        hidden = self.W_xh(_input).add(self.W_hh(hidden))
        hidden = self.act(hidden)
        output = self.W_hy(hidden)
        output = self.dropout(output)
        output = self.softmax(output)
        return output, hidden

    def initHidden(self):
        return torch.zeros(1, self.hidden_size)
```



Loss function: Negative Log Likelihood Loss



Optimization function: Adam

- Adam was presented by [Diederik Kingma](#) from OpenAI and [Jimmy Ba](#) from the University of Toronto in their 2015 [ICLR](#) paper (poster) titled "[Adam: A Method for Stochastic Optimization](#)". I will quote liberally from their paper in this post, unless stated otherwise.
- The authors describe Adam as combining the advantages of two other extensions of stochastic gradient descent. Specifically:
 - Adaptive Gradient Algorithm** (AdaGrad)
 - Root Mean Square Propagation** (RMSProp)

Training

```
loss_fun = nn.NLLLoss()
optimizer = torch.optim.Adam(rnn.parameters(), lr=0.01)

while epoch < 10:
    if p + seq_length + 1 >= len(data):
        epoch += 1
        p = 0
    # Generate input and target
    input_line_tensor, target_line_tensor = createTrainingExample(data[p:p+seq_length+1])
    target_line_tensor.unsqueeze_(-1)
    # Initiation
    hidden = rnn.initHidden()
    rnn.zero_grad()
    loss = 0
    # Training
    for i in range(input_line_tensor.size(0)):
        output, hidden = rnn(input_line_tensor[i], hidden)

        l = loss_fun(output, target_line_tensor[i])
        loss += l
    # Back propagation
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

    p += seq_length # move data pointer
```

Character-Level Predict Model

- Unsmoothed Maximum Likelihood Model
 - c is character, h is a n letters history
 - $P(c|h)$ stands for how likely is it to see c after we've seen h.
- Training data: Shakespeare
 - demo: A 4-letter model
 - Deterministic model
 - High dimensions

```
def train_char_lm(fname, order=4):
    data = file(fname).read()
    lm = defaultdict(Counter)
    pad = "~" * order
    data = pad + data
    for i in xrange(len(data)-order):
        history, char = data[i:i+order], data[i+order]
        lm[history][char] += 1
    def normalize(counter):
        s = float(sum(counter.values()))
        return [(c, cnt/s) for c, cnt in counter.iteritems()]
    outlm = {hist: normalize(chars) for hist, chars in lm.iteritems()}
    return outlm
```

```
lm['ello']
```

```
[('!', 0.0068143100511073255),
 (' ', 0.013628620102214651),
 ('"', 0.017035775127768313),
 ('.', 0.027257240204429302),
 (''', 0.0068143100511073255),
 ('r', 0.059625212947189095),
 ('u', 0.03747870528109029),
 ('w', 0.817717206132879),
 ('n', 0.0017035775127768314),
 (':', 0.005110732538330494),
 ('?', 0.0068143100511073255)]
```

Result from 10-letter model

First Citizen:
Nay, then, that was hers,
It speaks against your other service:
But since the
youth of the circumstance be spoken:
Your uncle and one Baptista's daughter.

SEBASTIAN:
Do I stand till the break off.

BIRON:
Hide thy head.

VENTIDIUS:
He purposeth to Athens: whither, with the vow
I made to handle you.

FALSTAFF:
My good knave.

MALVOLIO:
Sad, lady! I could be forgiven you, you're welcome. Give ear, sir, my doublet and hose and leave this present deat

Second Gentleman:
Who may that she confess it is my lord enraged and forestalled ere we come to be a man. Drown thyself?

Three layers RNN model

This is a 3-layer RNN with 512 hidden nodes on each layer.

We will need

1. 68 x 512
2. 512 x 512
3. 512 x 512
4. 512 x 512
5. 512 x 68

856,064 parameters in total

$N_{\text{para}} \times 4 \text{ bytes} = 3.42 \text{ MB}$

Backward pass requires 3 times of the memory

$3.42 \times 3 \sim 10.26 \text{ MB}$

```
class RNN3(nn.Module):
    def __init__(self, input_size=68, hidden_size=512, output_size=68, n_layers=3):
        super().__init__()
        self.hidden_size = hidden_size
        self.n_layers = n_layers
        self.rnn = nn.RNN(input_size, hidden_size, n_layers, dropout=0.1)
        self.fc = nn.Linear(hidden_size, output_size)
        self.softmax = nn.LogSoftmax(dim=2)

    def forward(self, input):
        """
        input: a batch of one hot charactors [25, 1, 68]"""
        hidden = self.initHidden(input.size(1))
        return_ = []
        for i in range(input.size(0)):
            _input = input[i:i+1]
            output, hidden = self.rnn(_input, hidden)
            output = self.fc(output)
            output = self.softmax(output)
            return_.append(output)
        return return_

    def prediction(self, input, hidden):
        """
        input: a batch of one hot charactors [1, 1, 68]"""
        output, hidden = self.rnn(input, hidden)
        output = self.fc(output)
        output = self.softmax(output)
        return output, hidden

    def initHidden(self, batch_size):
        return torch.zeros(self.n_layers, batch_size, self.hidden_size).cuda()
```

Three layers RNN model

PANDARUS:

Alas, I think he shall be come approached and the day
When little strain would be attain'd into being never fed,
And who is but a chain and subjects of his death,
I should not sleep.

Second Senator:

They are away this miseries, produced upon my soul,
Breaking and strongly should be buried, when I perish
The earth and thoughts of many states.

DUKE VINCENTIO:

Well, your wit is in the care of side and that.

Second Lord:

They would be ruled after this chamber, and
my fair nudes begun out of the fact, to be conveyed,
Whose noble souls I'll have the heart of the wars.

Clown:

Come, sir, I will make did behold your worship.

VIOLA:

I'll drink it.

For xml format (multi-layer LSTM)

```
<page>
  <title>Antichrist</title>
  <id>865</id>
  <revision>
    <id>15900676</id>
    <timestamp>2002-08-03T18:14:12Z</timestamp>
    <contributor>
      <username>Paris</username>
      <id>23</id>
    </contributor>
    <minor />
    <comment>Automated conversion</comment>
    <text xml:space="preserve">#REDIRECT [[Christianity]]</text>
  </revision>
</page>
```


Even Latex (multi-layer LSTM)

```
\begin{proof}
We may assume that  $\mathcal{I}$  is an abelian sheaf on  $\mathcal{C}$ .
\item Given a morphism  $\Delta : \mathcal{F} \rightarrow \mathcal{I}$ 
is an injective and let  $q$  be an abelian sheaf on  $X$ .
Let  $\mathcal{F}$  be a fibered complex. Let  $\mathcal{F}$  be a category.
\begin{enumerate}
\item \hyperref[setain-construction-phantom]{Lemma}
\label{lemma-characterize-quasi-finite}
Let  $\mathcal{F}$  be an abelian quasi-coherent sheaf on  $\mathcal{C}$ .
Let  $\mathcal{F}$  be a coherent  $\mathcal{O}_X$ -module. Then
 $\mathcal{F}$  is an abelian catenary over  $\mathcal{C}$ .
\item The following are equivalent
\begin{enumerate}
\item  $\mathcal{F}$  is an  $\mathcal{O}_X$ -module.
\end{enumerate}
\end{enumerate}
\end{lemma}
```

Thanks

- References

- <https://pytorch.org/tutorials/>
- <https://nbviewer.jupyter.org/gist/yoavg/d76121dfde2618422139>
- <https://karpathy.github.io/2015/05/21/rnn-effectiveness/>
- <https://cs.stanford.edu/people/karpathy/char-rnn/shakespear.txt>
- <https://gist.github.com/karpathy/d4dee566867f8291f086>

- Code on github

- https://github.com/bobaoai/Sample_RNN.git